

A Short History of Markov Chain Monte Carlo: Subjective Recollections from Incomplete Data¹

Christian Robert and George Casella

This paper is dedicated to the memory of our friend Julian Besag, a giant in the field of MCMC.

Abstract. We attempt to trace the history and development of Markov chain Monte Carlo (MCMC) from its early inception in the late 1940s through its use today. We see how the earlier stages of Monte Carlo (MC, not MCMC) research have led to the algorithms currently in use. More importantly, we see how the development of this methodology has not only changed our solutions to problems, but has changed the way we think about problems.

Key words and phrases: Gibbs sampling, Metropolis–Hasting algorithm, hierarchical models, Bayesian methods.

1. INTRODUCTION

Markov chain Monte Carlo (MCMC) methods have been around for almost as long as Monte Carlo techniques, even though their impact on Statistics has not been truly felt until the very early 1990s, except in the specialized fields of Spatial Statistics and Image Analysis, where those methods appeared earlier. The emergence of Markov based techniques in Physics is a story that remains untold within this survey (see Landau and Binder, 2005). Also, we will not enter into a description of MCMC techniques. A comprehensive treatment of MCMC techniques, with further references, can be found in Robert and Casella (2004).

Christian Robert is Professor of Statistics, CEREMADE, Université Paris Dauphine, 75785 Paris cedex 16, France (e-mail: xian@ceremade.dauphine.fr). George Casella is Distinguished Professor, Department of Statistics, University of Florida, Gainesville, Florida 32611, USA (e-mail: casella@ufl.edu).

¹A shorter version of this paper appears as Chapter 1 in *Handbook of Markov Chain Monte Carlo* (2011), edited by Steve Brooks, Andrew Gelman, Galin Jones and Xiao-Li Meng. Chapman & Hall/CRC Handbooks of Modern Statistical Methods, Boca Raton, Florida.

We will distinguish between the introduction of Metropolis–Hastings based algorithms and those related to Gibbs sampling, since they each stem from radically different origins, even though their mathematical justification via Markov chain theory is the same. Tracing the development of Monte Carlo methods, we will also briefly mention what we might call the “second-generation MCMC revolution.” Starting in the mid-to-late 1990s, this includes the development of particle filters, reversible jump and perfect sampling, and concludes with more current work on population or sequential Monte Carlo and regeneration and the computing of “honest” standard errors.

As mentioned above, the realization that Markov chains could be used in a wide variety of situations only came (to mainstream statisticians) with Gelfand and Smith (1990), despite earlier publications in the statistical literature like Hastings (1970), Geman and Geman (1984) and Tanner and Wong (1987). Several reasons can be advanced: lack of computing machinery (think of the computers of 1970!), or background on Markov chains, or hesitation to trust in the practicality of the method. It thus required visionary researchers like Gelfand and Smith to convince the community, supported by papers that demonstrated, through a series of applications, that the method was easy to un-

derstand, easy to implement and practical (Gelfand et al., 1990; Gelfand, Smith and Lee, 1992; Smith and Gelfand, 1992; Wakefield et al., 1994). The rapid emergence of the dedicated BUGS (Bayesian inference Using Gibbs Sampling) software as early as 1991, when a paper on BUGS was presented at the Valencia meeting, was another compelling argument for adopting, at large, MCMC algorithms.²

2. BEFORE THE REVOLUTION

Monte Carlo methods were born in Los Alamos, New Mexico during World War II, eventually resulting in the Metropolis algorithm in the early 1950s. While Monte Carlo methods were in use by that time, MCMC was brought closer to statistical practicality by the work of Hastings in the 1970s.

What can be reasonably seen as the first MCMC algorithm is what we now call the Metropolis algorithm, published by Metropolis et al. (1953). It emanates from the same group of scientists who produced the Monte Carlo method, namely, the research scientists of Los Alamos, mostly physicists working on mathematical physics and the atomic bomb.

MCMC algorithms therefore date back to the same time as the development of regular (MC only) Monte Carlo methods, which are usually traced to Ulam and von Neumann in the late 1940s. Stanislaw Ulam associates the original idea with an intractable combinatorial computation he attempted in 1946 (calculating the probability of winning at the card game “solitaire”). This idea was enthusiastically adopted by John von Neumann for implementation with direct applications to neutron diffusion, the name “Monte Carlo” being suggested by Nicholas Metropolis. (Eckhardt, 1987, describes these early Monte Carlo developments, and Hitchcock, 2003, gives a brief history of the Metropolis algorithm.)

These occurrences very closely coincide with the appearance of the very first computer, the ENIAC, which came to life in February 1946, after three years of construction. The Monte Carlo method was set up by von Neumann, who was using it on thermonuclear and fission problems as early as 1947. At the same time, that is, 1947, Ulam and von Neumann invented inversion and accept-reject techniques (also recounted in

Eckhardt, 1987) to simulate from nonuniform distributions. Without computers, a rudimentary version invented by Fermi in the 1930s did not get any recognition (Metropolis, 1987). Note also that, as early as 1949, a symposium on Monte Carlo was supported by Rand, NBS and the Oak Ridge laboratory and that Metropolis and Ulam (1949) published the very first paper about the Monte Carlo method.

2.1 The Metropolis et al. (1953) Paper

The first MCMC algorithm is associated with a second computer, called MANIAC, built³ in Los Alamos under the direction of Metropolis in early 1952. Both a physicist and a mathematician, Nicolas Metropolis, who died in Los Alamos in 1999, came to this place in April 1943. The other members of the team also came to Los Alamos during those years, including the controversial Teller. As early as 1942, he became obsessed with the hydrogen (H) bomb, which he eventually managed to design with Stanislaw Ulam, using the better computer facilities in the early 1950s.

Published in June 1953 in the *Journal of Chemical Physics*, the primary focus of Metropolis et al. (1953) is the computation of integrals of the form

$$\mathfrak{J} = \frac{\int F(\theta) \exp\{-E(\theta)/kT\} d\theta}{\int \exp\{-E(\theta)/kT\} d\theta},$$

on \mathbb{R}^{2N} , θ denoting a set of N particles on \mathbb{R}^2 , with the energy E being defined as

$$E(\theta) = \frac{1}{2} \sum_{i=1}^N \sum_{\substack{j=1 \\ j \neq i}}^N V(d_{ij}),$$

where V is a potential function and d_{ij} the Euclidean distance between particles i and j in θ . The *Boltzmann distribution* $\exp\{-E(\theta)/kT\}$ is parameterized by the *temperature* T , k being the Boltzmann constant, with a normalization factor

$$Z(T) = \int \exp\{-E(\theta)/kT\} d\theta,$$

that is not available in closed form, except in trivial cases. Since θ is a $2N$ -dimensional vector, numerical integration is impossible. Given the large dimension of the problem, even standard Monte Carlo techniques fail to correctly approximate \mathfrak{J} , since $\exp\{-E(\theta)/kT\}$

²Historically speaking, the development of BUGS initiated from Geman and Geman (1984) and Pearl (1987), in accord with the developments in the artificial intelligence community, and it predates Gelfand and Smith (1990).

³MANIAC stands for *Mathematical Analyzer, Numerical Integrator and Computer*.

is very small for most realizations of the random configurations of the particle system (uniformly in the $2N$ square). In order to improve the efficiency of the Monte Carlo method, Metropolis et al. (1953) propose a random walk modification of the N particles. That is, for each particle i ($1 \leq i \leq N$), values

$$x'_i = x_i + \sigma \xi_{1i} \quad \text{and} \quad y'_i = y_i + \sigma \xi_{2i}$$

are proposed, where both ξ_{1i} and ξ_{2i} are uniform $\mathcal{U}(-1, 1)$. The energy difference ΔE between the new configuration and the previous one is then computed and the new configuration is accepted with probability

$$(1) \quad \min\{1, \exp(-\Delta E/kT)\},$$

and otherwise the previous configuration is replicated, in the sense that its counter is increased by one in the final average of the $F(\theta_t)$'s over the τ moves of the random walk, $1 \leq t \leq \tau$. Note that Metropolis et al. (1953) move one particle at a time, rather than moving all of them together, which makes the initial algorithm appear as a primitive kind of Gibbs sampler!

The authors of Metropolis et al. (1953) demonstrate the validity of the algorithm by first establishing irreducibility, which they call *ergodicity*, and second proving ergodicity, that is, convergence to the stationary distribution. The second part is obtained via a discretization of the space: They first note that the proposal move is reversible, then establish that $\exp\{-E/kT\}$ is invariant. The result is therefore proven in its full generality, minus the discretization. The number of iterations of the Metropolis algorithm used in the paper seems to be limited: 16 steps for burn-in and 48 to 64 subsequent iterations, which required four to five hours on the Los Alamos computer.

An interesting variation is the *Simulated Annealing* algorithm, developed by Kirkpatrick, Gelatt and Vecchi (1983), who connected optimization with *annealing*, the cooling of a metal. Their variation is to allow the temperature T in (1) to change as the algorithm runs, according to a “cooling schedule,” and the Simulated Annealing algorithm can be shown to find the global maximum with probability 1, although the analysis is quite complex due to the fact that, with varying T , the algorithm is no longer a time-homogeneous Markov chain.

2.2 The Hastings (1970) Paper

The Metropolis algorithm was later generalized by Hastings (1970) and his student Peskun (1973, 1981) as a statistical simulation tool that could overcome the curse of dimensionality met by regular Monte Carlo

methods, a point already emphasized in Metropolis et al. (1953).⁴

In his *Biometrika* paper,⁵ Hastings (1970) also defines his methodology for finite and reversible Markov chains, treating the continuous case by using a discretization analogy. The generic probability of acceptance for a move from state i to state j is

$$\alpha_{ij} = \frac{s_{ij}}{1 + (\pi_i/\pi_j)(q_{ij}/q_{ji})},$$

where $s_{ij} = s_{ji}$, π_i denotes the target and q_{ij} the proposal. This generic form of probability encompasses the forms of both Metropolis et al. (1953) and Barker (1965). At this stage, Hastings mentions that *little is known about the relative merits of those two choices* (even though) *Metropolis's method may be preferable*. He also warns against *high rejection rates as indicative of a poor choice of transition matrix*, but does not mention the opposite pitfall of low rejection rates, associated with a slow exploration of the target.

The examples given in the paper are a Poisson target with a ± 1 random walk proposal, a normal target with a uniform random walk proposal mixed with its reflection, that is, a uniform proposal centered at $-\theta_t$ rather than at the current value θ_t of the Markov chain, and then a multivariate target where Hastings introduces a Gibbs sampling strategy, updating one component at a time and defining the composed transition as satisfying the stationary condition because each component does leave the target invariant. Hastings (1970) actually refers to Ehrman, Fosdick and Handscomb (1960) as a preliminary, if specific, instance of this sampler. More precisely, this is Metropolis-within-Gibbs except for the name. This first introduction of the Gibbs sampler has thus been completely overlooked, even though the proof of convergence is completely general, based on a composition argument as in Tierney (1994), discussed in Section 4.1. The remainder of the paper deals with (a) an importance sampling version of MCMC, (b) general remarks about assessment of the error, and (c) an application to random orthogonal matrices, with another example of Gibbs sampling.

Three years later, Peskun (1973) published a comparison of Metropolis' and Barker's forms of acceptance probabilities and shows in a discrete setup that

⁴In fact, Hastings starts by mentioning a decomposition of the target distribution into a *product of one-dimensional conditional distributions*, but this falls short of an early Gibbs sampler.

⁵Hastings (1970) is one of the ten *Biometrika* papers reproduced in Titterton and Cox (2001).

the optimal choice is that of Metropolis, where optimality is to be understood in terms of the asymptotic variance of any empirical average. The proof is a direct consequence of a result by Kemeny and Snell (1960) on the asymptotic variance. Peskun also establishes that this asymptotic variance can improve upon the i.i.d. case if and only if the eigenvalues of $\mathbf{P} - \mathbf{A}$ are all negative, when \mathbf{A} is the transition matrix corresponding to i.i.d. simulation and \mathbf{P} the transition matrix corresponding to the Metropolis algorithm, but he concludes that the trace of $\mathbf{P} - \mathbf{A}$ is always positive.

3. SEEDS OF THE REVOLUTION

A number of earlier pioneers had brought forward the seeds of Gibbs sampling; in particular, Hammersley and Clifford had produced a constructive argument in 1970 to recover a joint distribution from its conditionals, a result later called the *Hammersley–Clifford* theorem by Besag (1974, 1986). Besides Hastings (1970) and Geman and Geman (1984), already mentioned, other papers that contained the seeds of Gibbs sampling are Besag and Clifford (1989), Broniatowski, Celeux and Diebolt (1984), Qian and Titterton (1990) and Tanner and Wong (1987).

3.1 Besag’s Early Work and the Fundamental (Missing) Theorem

In the early 1970’s, Hammersley, Clifford and Besag were working on the specification of joint distributions from conditional distributions and on necessary and sufficient conditions for the conditional distributions to be compatible with a joint distribution. What is now known as the *Hammersley–Clifford* theorem states that a joint distribution for a vector associated with a dependence graph (edge meaning dependence and absence of edge conditional independence) must be represented as a product of functions over the *cliques* of the graphs, that is, of functions depending only on the components indexed by the labels in the clique.⁶

From a historical point of view, Hammersley (1974) explains why the Hammersley–Clifford theorem was never published as such, but only through Besag (1974). The reason is that Clifford and Hammersley were dissatisfied with the positivity constraint: The joint density could be recovered from the full conditionals only when the support of the joint was made

of the product of the supports of the full conditionals. While they strived to *make the theorem independent of any positivity condition*, their graduate student published a counter-example that put a full stop to their endeavors (Moussouris, 1974).

While Besag (1974) can certainly be credited to some extent of the (re-)discovery of the Gibbs sampler, Besag (1975) expressed doubt about the practicality of his method, noting that “the technique is unlikely to be particularly helpful in many other than binary situations and the Markov chain itself has no practical interpretation,” clearly understating the importance of his own work.

A more optimistic sentiment was expressed earlier by Hammersley and Handscomb (1964) in their textbook on Monte Carlo methods. There they cover such topics as “Crude Monte Carlo,” importance sampling, control variates and “Conditional Monte Carlo,” which looks surprisingly like a missing-data completion approach. Of course, they do not cover the Hammersley–Clifford theorem, but they state in the Preface:

We are convinced nevertheless that Monte Carlo methods will one day reach an impressive maturity.

Well said!

3.2 EM and Its Simulated Versions as Precursors

Because of its use for missing data problems, the EM algorithm (Dempster, Laird and Rubin, 1977) has early connections with Gibbs sampling. For instance, Broniatowski, Celeux and Diebolt (1984) and Celeux and Diebolt (1985) had tried to overcome the dependence of EM methods on the starting value by replacing the E step with a *simulation* step, the missing data z being generated conditionally on the observation x and on the current value of the parameter θ_m . The maximization in the M step is then done on the simulated complete-data log-likelihood, a predecessor to the Gibbs step of Diebolt and Robert (1994) for mixture estimation. Unfortunately, the theoretical convergence results for these methods are limited. Celeux and Diebolt (1990) have, however, solved the convergence problem of SEM by devising a hybrid version called SAEM (for *Simulated Annealing EM*), where the amount of randomness in the simulations decreases with the iterations, ending up with an EM algorithm.⁷

⁶A clique is a maximal subset of the nodes of a graphs such that every pair of nodes within the clique is connected by an edge (Cressie, 1993).

⁷Other and more well-known connections between EM and MCMC algorithms can be found in the literature (Liu and Rubin,

3.3 Gibbs and Beyond

Although somewhat removed from statistical inference in the classical sense and based on earlier techniques used in Statistical Physics, the landmark paper by Geman and Geman (1984) brought Gibbs sampling into the arena of statistical application. This paper is also responsible for the name *Gibbs sampling*, because it implemented this method for the Bayesian study of *Gibbs random fields* which, in turn, derive their name from the physicist Josiah Willard Gibbs (1839–1903). This original implementation of the Gibbs sampler was applied to a discrete image processing problem and did not involve completion. But this was one more spark that led to the explosion, as it had a clear influence on Green, Smith, Spiegelhalter and others.

The extent to which Gibbs sampling and Metropolis algorithms were in use within the image analysis and point process communities is actually quite large, as illustrated in Ripley (1987) where Section 4.7 is entitled “Metropolis’ method and random fields” and describes the implementation and the validation of the Metropolis algorithm in a finite setting with an application to Markov random fields and the corresponding issue of bypassing the normalizing constant. Besag, York and Mollié (1991) is another striking example of the activity in the spatial statistics community at the end of the 1980s.

4. THE REVOLUTION

The gap of more than 30 years between Metropolis et al. (1953) and Gelfand and Smith (1990) can still be partially attributed to the lack of appropriate computing power, as most of the examples now processed by MCMC algorithms could not have been treated previously, even though the hundreds of dimensions processed in Metropolis et al. (1953) were quite formidable. However, by the mid-1980s, the pieces were all in place.

After Peskun, MCMC in the statistical world was dormant for about 10 years, and then several papers appeared that highlighted its usefulness in specific settings like pattern recognition, image analysis or spatial statistics. In particular, Geman and Geman (1984) influenced Gelfand and Smith (1990) to write a paper that is the genuine starting point for an intensive use of

MCMC methods by the mainstream statistical community. It sparked new interest in Bayesian methods, statistical computing, algorithms and stochastic processes through the use of computing algorithms such as the Gibbs sampler and the Metropolis–Hastings algorithm. (See Casella and George, 1992, for an elementary introduction to the Gibbs sampler.⁸)

Interestingly, the earlier paper by Tanner and Wong (1987) had essentially the same ingredients as Gelfand and Smith (1990), namely, the fact that simulating from the conditional distributions is sufficient to asymptotically simulate from the joint. This paper was considered important enough to be a discussion paper in the *Journal of the American Statistical Association*, but its impact was somehow limited, compared with Gelfand and Smith (1990). There are several reasons for this; one being that the method seemed to only apply to missing data problems, this impression being reinforced by the name *data augmentation*, and another is that the authors were more focused on approximating the posterior distribution. They suggested a MCMC approximation to the target $\pi(\theta|x)$ at each iteration of the sampler, based on

$$\frac{1}{m} \sum_{k=1}^m \pi(\theta|x, z^{t,k}),$$

$$z^{t,k} \sim \hat{\pi}_{t-1}(z|x), \quad k = 1, \dots, m,$$

that is, by replicating m times the simulations from the current approximation $\hat{\pi}_{t-1}(z|x)$ of the marginal posterior distribution of the missing data. This focus on estimation of the posterior distribution connected the original Data Augmentation algorithm to EM, as pointed out by Dempster in the discussion. Although the discussion by Carl Morris gets very close to the two-stage Gibbs sampler for hierarchical models, he is still concerned about doing m iterations, and worries about how costly that would be. Tanner and Wong mention taking $m = 1$ at the end of the paper, referring to this as an “extreme case.”

In a sense, Tanner and Wong (1987) were still too close to Rubin’s 1978 multiple imputation to start a new revolution. Yet another reason for this may be

1994; Meng and Rubin, 1993; Wei and Tanner, 1990), but the connection with Gibbs sampling is more tenuous in that the simulation methods are used to approximate quantities in a Monte Carlo fashion.

⁸On a humorous note, the original Technical Report of this paper was called *Gibbs for Kids*, which was changed because a referee did not appreciate the humor. However, our colleague Dan Gianola, an Animal Breeder at Wisconsin, liked the title. In using Gibbs sampling in his work, he gave a presentation in 1993 at the 44th Annual Meeting of the European Association for Animal Production, Aarhus, Denmark. The title: *Gibbs for Pigs*.

that the theoretical background was based on functional analysis rather than Markov chain theory, which needed, in particular, for the Markov kernel to be uniformly bounded and equicontinuous. This may have discouraged potential users as requiring too much mathematics.

The authors of this review were fortunate enough to attend many focused conferences during this time, where we were able to witness the explosion of Gibbs sampling. In the summer of 1986 in Bowling Green, Ohio, Adrian Smith gave a series of ten lectures on hierarchical models. Although there was a lot of computing mentioned, the Gibbs sampler was not fully developed yet. In another lecture in June 1989 at a Bayesian workshop in Sherbrooke, Québec, he revealed for the first time the generic features of the Gibbs sampler, and we still remember vividly the shock induced on ourselves and on the whole audience by the sheer breadth of the method: This development of Gibbs sampling, MCMC, and the resulting seminal paper of Gelfand and Smith (1990) was an *epiphany* in the world of Statistics.

DEFINITION (Epiphany n). A spiritual event in which the essence of a given object of manifestation appears to the subject, as in a sudden flash of recognition.

The explosion had begun, and just two years later, at an MCMC conference at Ohio State University organized by Alan Gelfand, Prem Goel and Adrian Smith, there were three full days of talks. The presenters at the conference read like a Who's Who of MCMC, and the level, intensity and impact of that conference, and the subsequent research, are immeasurable. Many of the talks were to become influential papers, including Albert and Chib (1993), Gelman and Rubin (1992), Geyer (1992), Gilks (1992), Liu, Wong and Kong (1994, 1995) and Tierney (1994). The program of the conference is reproduced in the [Appendix](#).

Approximately one year later, in May of 1992, there was a meeting of the Royal Statistical Society on "The Gibbs sampler and other Markov chain Monte Carlo methods," where four papers were presented followed by much discussion. The papers appear in the first volume of *JRSSB* in 1993, together with 49(!) pages of discussion. The excitement is clearly evident in the writings, even though the theory and implementation were not always perfectly understood.⁹

⁹On another humorous note, Peter Clifford opened the discussion by noting "...we have had the opportunity to hear a large amount

Looking at these meetings, we can see the paths that Gibbs sampling would lead us down. In the next two sections we will summarize some of the advances from the early to mid 1990s.

4.1 Advances in MCMC Theory

Perhaps the most influential MCMC theory paper of the 1990s is Tierney (1994), who carefully laid out all of the assumptions needed to analyze the Markov chains and then developed their properties, in particular, convergence of ergodic averages and central limit theorems. In one of the discussions of that paper, Chan and Geyer (1994) were able to relax a condition on Tierney's Central Limit Theorem, and this new condition plays an important role in research today (see Section 5.4). A pair of very influential, and innovative, papers is the work of Liu, Wong and Kong (1994, 1995), who very carefully analyzed the covariance structure of Gibbs sampling, and were able to formally establish the validity of Rao–Blackwellization in Gibbs sampling. Gelfand and Smith (1990) had used Rao–Blackwellization, but it was not justified at that time, as the original theorem was only applicable to i.i.d. sampling, which is not the case in MCMC. Another significant entry is Rosenthal (1995), who obtained one of the earliest results on exact rates of convergence.

Another paper must be singled out, namely, Mengersen and Tweedie (1996), for setting the tone for the study of the speed of convergence of MCMC algorithms to the target distribution. Subsequent works in this area by Richard Tweedie, Gareth Roberts, Jeff Rosenthal and co-authors are too numerous to be mentioned here, even though the paper by Roberts, Gelman and Gilks (1997) must be cited for setting explicit targets on the acceptance rate of the random walk Metropolis–Hastings algorithm, as well as Roberts and Rosenthal (1999) for getting an upper bound on the number of iterations (523) needed to approximate the target up to 1% by a slice sampler. The untimely death of Richard Tweedie in 2001, alas, had a major impact on the book about MCMC convergence he was contemplating with Gareth Roberts.

about an important new area in statistics. It may well be remembered as the 'afternoon of the 11 Bayesians.' Bayesianism has obviously come a long way. It used to be that you could tell a Bayesian by his tendency to hold meetings in isolated parts of Spain and his obsession with coherence, self-interrogation and other manifestations of paranoia. Things have changed, and there may be a general lesson here for statistics. Isolation is counter-productive."

One pitfall arising from the widespread use of Gibbs sampling was the tendency to specify models only through their conditional distributions, almost always without referring to the positivity conditions in Section 3. Unfortunately, it is possible to specify a perfectly legitimate-looking set of conditionals that do not correspond to any joint distribution, and the resulting Gibbs chain cannot converge. Hobert and Casella (1996) were able to document the conditions needed for a convergent Gibbs chain, and alerted the Gibbs community to this problem, which only arises when improper priors are used, but this is a frequent occurrence.

Much other work followed, and continues to grow today. Geyer and Thompson (1995) describe how to put a “ladder” of chains together to have both “hot” and “cold” exploration, followed by Neal’s 1996 introduction of tempering; Athreya, Doss and Sethuraman (1996) gave more easily verifiable conditions for convergence; Meng and van Dyk (1999) and Liu and Wu (1999) developed the theory of parameter expansion in the Data Augmentation algorithm, leading to construction of chains with faster convergence, and to the work of Hobert and Marchev (2008), who give precise constructions and theorems to show how parameter expansion can uniformly improve over the original chain.

4.2 Advances in MCMC Applications

The real reason for the explosion of MCMC methods was the fact that an enormous number of problems that were deemed to be computational nightmares now cracked open like eggs. As an example, consider this very simple random effects model from Gelfand and Smith (1990). Observe

$$(2) \quad Y_{ij} = \theta_i + \varepsilon_{ij}, \quad i = 1, \dots, K, \quad j = 1, \dots, J,$$

where

$$\theta_i \sim N(\mu, \sigma_\theta^2),$$

$$\varepsilon_{ij} \sim N(0, \sigma_\varepsilon^2), \quad \text{independent of } \theta_i.$$

Estimation of the variance components can be difficult for a frequentist (REML is typically preferred), but it indeed was a nightmare for a Bayesian, as the integrals were intractable. However, with the usual priors on μ , σ_θ^2 and σ_ε^2 , the full conditionals are trivial to sample from and the problem is easily solved via Gibbs sampling. Moreover, we can increase the number of variance components and the Gibbs solution remains easy to implement.

During the early 1990s, researchers found that Gibbs, or Metropolis–Hastings, algorithms would be able to give solutions to almost any problem that they looked at, and there was a veritable flood of papers applying MCMC to previously intractable models, and getting good answers. For example, building on (2), it was quickly realized that Gibbs sampling was an easy route to getting estimates in the linear mixed models (Wang, Rutledge and Gianola, 1993, 1994), and even generalized linear mixed models (Zeger and Karim, 1991). Building on the experience gained with the EM algorithm, similar arguments made it possible to analyze probit models using a latent variable approach in a linear mixed model (Albert and Chib, 1993), and in mixture models with Gibbs sampling (Diebolt and Robert, 1994). It progressively dawned on the community that latent variables could be artificially introduced to run the Gibbs sampler in about every situation, as eventually published in Damien, Wakefield and Walker (1999), the main example being the slice sampler (Neal, 2003). A very incomplete list of some other applications include changepoint analysis (Carlin, Gelfand and Smith, 1992; Stephens, 1994), Genomics (Churchill, 1995; Lawrence et al., 1993; Stephens and Smith, 1993), capture–recapture (Dupuis, 1995; George and Robert, 1992), variable selection in regression (George and McCulloch, 1993), spatial statistics (Raftery and Banfield, 1991), and longitudinal studies (Lange, Carlin and Gelfand, 1992).

Many of these applications were advanced through other developments such as the Adaptive Rejection Sampling of Gilks (1992); Gilks, Best and Tan (1995), and the simulated tempering approaches of Geyer and Thompson (1995) or Neal (1996).

5. AFTER THE REVOLUTION

After the revolution comes the “second” revolution, but now we have a more mature field. The revolution has slowed, and the problems are being solved in, perhaps, deeper and more sophisticated ways, even though Gibbs sampling also offers to the amateur the possibility to handle Bayesian analysis in complex models at little cost, as exhibited by the widespread use of BUGS, which mostly focuses¹⁰ on this approach. But, as before, the methodology continues to expand the set of problems for which statisticians can provide meaningful solutions, and thus continues to further the impact of Statistics.

¹⁰BUGS now uses both Gibbs sampling and Metropolis–Hastings algorithms.

5.1 A Brief Glimpse at Particle Systems

The realization of the possibilities of iterating importance sampling is not new: in fact, it is about as old as Monte Carlo methods themselves. It can be found in the molecular simulation literature of the 50s, as in Hammersley and Morton (1954), Rosenbluth and Rosenbluth (1955) and Marshall (1965). Hammersley and colleagues proposed such a method to simulate a self-avoiding random walk (see Madras and Slade, 1993) on a grid, due to huge inefficiencies in regular importance sampling and rejection techniques. Although this early implementation occurred in particle physics, the use of the term “particle” only dates back to Kitagawa (1996), while Carpenter, Clifford and Fernhead (1997) coined the term “particle filter.” In signal processing, early occurrences of a particle filter can be traced back to Handschin and Mayne (1969).

More in connection with our theme, the landmark paper of Gordon, Salmond and Smith (1993) introduced the bootstrap filter which, while formally connected with importance sampling, involves past simulations and possible MCMC steps (Gilks and Berzuini, 2001). As described in the volume edited by Doucet, de Freitas and Gordon (2001), particle filters are simulation methods adapted to sequential settings where data are collected progressively in time, as in radar detection, telecommunication correction or financial volatility estimation. Taking advantage of state-space representations of those dynamic models, particle filter methods produce Monte Carlo approximations to the posterior distributions by propagating simulated samples whose weights are actualized against the incoming observations. Since the importance weights have a tendency to degenerate, that is, all weights but one are close to zero, additional MCMC steps can be introduced at times to recover the variety and representativeness of the sample. Modern connections with MCMC in the construction of the proposal kernel are to be found, for instance, in Doucet, Godsill and Andrieu (2000) and in Del Moral, Doucet and Jasra (2006). At the same time, sequential imputation was developed in Kong, Liu and Wong (1994), while Liu and Chen (1995) first formally pointed out the importance of re-sampling in sequential Monte Carlo, a term coined by them.

The recent literature on the topic more closely bridges the gap between sequential Monte Carlo and MCMC methods by making adaptive MCMC a possibility (see, e.g., Andrieu et al., 2004, or Roberts and Rosenthal, 2007).

5.2 Perfect Sampling

Introduced in the seminal paper of Propp and Wilson (1996), perfect sampling, namely, the ability to use MCMC methods to produce an exact (or perfect) simulation from the target, maintains a unique place in the history of MCMC methods. Although this exciting discovery led to an outburst of papers, in particular, in the large body of work of Møller and coauthors, including the book by Møller and Waagepetersen (2003), as well as many reviews and introductory materials, like Casella, Lavine and Robert (2001), Fisman (1998) and Dimakos (2001), the excitement quickly dried out. The major reason for this ephemeral lifespan is that the construction of perfect samplers is most often close to impossible or impractical, despite some advances in the implementation (Fill, 1998a, 1998b).

There is, however, ongoing activity in the area of point processes and stochastic geometry, much from the work of Møller and Kendall. In particular, Kendall and Møller (2000) developed an alternative to the *Coupling From The Past* (CFPT) algorithm of Propp and Wilson (1996), called *horizontal CFTP*, which mainly applies to point processes and is based on continuous time birth-and-death processes. See also Fernández, Ferrari and Garcia (1999) for another horizontal CFTP algorithm for point processes. Berthelsen and Møller (2003) exhibited a use of these algorithms for nonparametric Bayesian inference on point processes.

5.3 Reversible Jump and Variable Dimensions

From many viewpoints, the invention of the reversible jump algorithm in Green (1995) can be seen as the start of the second MCMC revolution: the formalization of a Markov chain that moves across models and parameter spaces allowed for the Bayesian processing of a wide variety of new models and contributed to the success of Bayesian model choice and subsequently to its adoption in other fields. There exist earlier alternative Monte Carlo solutions like Gelfand and Dey (1994) and Carlin and Chib (1995), the later being very close in spirit to reversible jump MCMC (as shown by the completion scheme of Brooks, Giudici and Roberts, 2003), but the definition of a proper balance condition on cross-model Markov kernels in Green (1995) gives a generic setup for exploring variable dimension spaces, even when the number of models under comparison is infinite. The impact of this new idea was clearly perceived when looking at the First European Conference on Highly Structured Stochastic Systems that took place in Rebild, Denmark,

the next year, organized by Stephen Lauritzen and Jesper Møller: a large majority of the talks were aimed at direct implementations of RJMCMC to various inference problems. The application of RJMCMC to mixture order estimation in the discussion paper of Richardson and Green (1997) ensured further dissemination of the technique. Continuing to develop RJMCMCt, Stephens (2000) proposed a continuous time version of RJMCMC, based on earlier ideas of Geyer and Møller (1994), but with similar properties (Cappé, Robert and Rydén, 2003), while Brooks, Giudici and Roberts (2003) made proposals for increasing the efficiency of the moves. In retrospect, while reversible jump is somehow unavoidable in the processing of very large numbers of models under comparison, as, for instance, in variable selection (Marin and Robert, 2007), the implementation of a complex algorithm like RJMCMC for the comparison of a few models is somewhat of an overkill since there may exist alternative solutions based on model specific MCMC chains, for example (Chen, Shao and Ibrahim, 2000).

5.4 Regeneration and the CLT

While the Central Limit Theorem (CLT) is a central tool in Monte Carlo convergence assessment, its use in MCMC setups took longer to emerge, despite early signals by Geyer (1992), and it is only recently that sufficiently clear conditions emerged. We recall that the Ergodic Theorem (see, e.g., Robert and Casella, 2004, Theorem 6.63) states that, if $(\theta_t)_t$ is a Markov chain with stationary distribution π , and $h(\cdot)$ is a function with finite variance, then under fairly mild conditions,

$$(3) \quad \lim_{n \rightarrow \infty} \bar{h}_n = \int h(\theta) \pi(\theta) d\theta = E_\pi h(\theta),$$

almost everywhere, where $\bar{h}_n = (1/n) \sum_{i=1}^n h(\theta_i)$. For the CLT to be used to monitor this convergence,

$$(4) \quad \frac{\sqrt{n}(\bar{h}_n - E_\pi h(\theta))}{\sqrt{\text{Var } h(\theta)}} \rightarrow N(0, 1),$$

there are two roadblocks. First, convergence to normality is strongly affected by the lack of independence. To get CLTs for Markov chains, we can use a result of Kipnis and Varadhan (1986), which requires the chain to be reversible, as is the case for holds for Metropolis–Hastings chains, or we must delve into mixing conditions (Billingsley, 1995, Section 27), which are typically not easy to verify. However, Chan and Geyer (1994) showed how the condition of geometric ergodicity could be used to establish CLTs for Markov chains. But getting the convergence is only half of the

problem. In order to use (4), we must be able to consistently estimate the variance, which turns out to be another difficult endeavor. The “naïve” estimate of the usual standard error is not consistent in the dependent case and the most promising paths for consistent variance estimates seems to be through regeneration and batch means.

The theory of regeneration uses the concept of a split chain (Athreya and Ney, 1978), and allows us to independently restart the chain while preserving the stationary distribution. These independent “tours” then allow the calculation of consistent variance estimates and honest monitoring of convergence through (4). Early work on applying regeneration to MCMC chains was done by Mykland, Tierney and Yu (1995) and Robert (1995), who showed how to construct the chains and use them for variance calculations and diagnostics (see also Guihenneuc-Jouyaux and Robert, 1998), as well as deriving adaptive MCMC algorithms (Gilks, Roberts and Sahu, 1998). Rosenthal (1995) also showed how to construct and use regenerative chains, and much of this work is reviewed in Jones and Hobert (2001). The most interesting and practical developments, however, are in Hobert et al. (2002) and Jones et al. (2006), where consistent estimators are constructed for $\text{Var } h(X)$, allowing valid monitoring of convergence in chains that satisfy the CLT. Interestingly, although Hobert et al. (2002) use regeneration, Jones et al. (2006) get their consistent estimators thorough another technique, that of cumulative batch means.

6. CONCLUSION

The impact of Gibbs sampling and MCMC was to change our entire method of thinking and attacking problems, representing a *paradigm shift* (Kuhn, 1996). Now, the collection of real problems that we could solve grew almost without bound. Markov chain Monte Carlo changed our emphasis from “closed form” solutions to algorithms, expanded our impact to solving “real” applied problems and to improving numerical algorithms using statistical ideas, and led us into a world where “exact” now means “simulated.”

This has truly been a quantum leap in the evolution of the field of statistics, and the evidence is that there are no signs of slowing down. Although the “explosion” is over, the current work is going deeper into theory and applications, and continues to expand our horizons and influence by increasing our ability to solve even bigger and more important problems. The size

of the data sets, and of the models, for example, in genomics or climatology, is something that could not have been conceived 60 years ago, when Ulam and von Neumann invented the Monte Carlo method. Now we continue to plod on, and hope that the advances that we make here will, in some way, help our colleagues 60 years in the future solve the problems that we cannot yet conceive.

APPENDIX: WORKSHOP ON BAYESIAN COMPUTATION

This section contains the program of the Workshop on *Bayesian Computation via Stochastic Simulation*, held at Ohio State University, February 15–17, 1991. The organizers, and their affiliations at the time, were Alan Gelfand, University of Connecticut, Prem Goel, Ohio State University, and Adrian Smith, Imperial College, London.

• Friday, Feb. 15, 1991.

- (a) Theoretical Aspect of Iterative Sampling, Chair: Adrian Smith.

(1) Martin Tanner, University of Rochester: *EM, MCEM, DA and PMDA*.

(2) Nick Polson, Carnegie Mellon University: *On the Convergence of the Gibbs Sampler and Its Rate*.

(3) Wing-Hung Wong, Augustin Kong and Jun Liu, University of Chicago: *Correlation Structure and Convergence of the Gibbs Sampler and Related Algorithms*.

- (b) Applications—I, Chair: Prem Goel.

(1) Nick Lange, Brown University, Brad Carlin, Carnegie Mellon University and Alan Gelfand, University of Connecticut: *Hierarchical Bayes Models for Progression of HIV Infection*.

(2) Cliff Litton, Nottingham University, England: *Archaeological Applications of Gibbs Sampling*.

(3) Jonas Mockus, Lithuanian Academy of Sciences, Vilnius: *Bayesian Approach to Global and Stochastic Optimization*.

• Saturday, Feb. 16, 1991.

- (a) Posterior Simulation and Markov Sampling, Chair: Alan Gelfand.

(1) Luke Tierney, University of Minnesota: *Exploring Posterior Distributions Using Markov Chains*.

(2) Peter Mueller, Purdue University: *A Generic Approach to Posterior Integration and Bayesian Sampling*.

(3) Andrew Gelman, University of California, Berkeley and Donald P. Rubin, Harvard University: *On the Routine Use of Markov Chains for Simulations*.

(4) Jon Wakefield, Imperial College, London: *Parameterization Issues in Gibbs Sampling*.

(5) Panickos Palettas, Virginia Polytechnic Institute: *Acceptance–Rejection Method in Posterior Computations*.

- (b) Applications—II, Chair: Mark Berliner.

(1) David Stephens, Imperial College, London: *Gene Mapping via Gibbs Sampling*.

(2) Constantine Gatsonis, Harvard University: *Random Effects Model for Ordinal Categorical Data with an Application to ROC Analysis*.

(3) Arnold Zellner, University of Chicago, Luc Bauwens and Herman Van Dijk: *Bayesian Specification Analysis and Estimation of Simultaneous Equation Models Using Monte Carlo Methods*.

- (c) Adaptive Sampling, Chair: Carl Morris.

(1) Mike Evans, University of Toronto and Carnegie Mellon University: *Some Uses of Adaptive Importance Sampling and Chaining*.

(2) Wally Gilks, Medical Research Council, Cambridge, England: *Adaptive Rejection Sampling*.

(3) Mike West, Duke University: *Mixture Model Approximations, Sequential Updating and Dynamic Models*.

• Sunday, Feb. 17, 1991.

- (a) Generalized Linear and Nonlinear Models, Chair: Rob Kass.

(1) Ruey Tsay and Robert McCulloch, University of Chicago: *Bayesian Analysis of Autoregressive Time Series*.

(2) Christian Ritter, University of Wisconsin: *Sampling Based Inference in Non Linear Regression*.

(3) William DuMouchel, BBN Software, Boston: *Application of the Gibbs Sampler to Variance Component Modeling*.

- (4) James Albert, Bowling Green University and Sidhartha Chib, Washington University, St. Louis: *Bayesian Regression Analysis of Binary Data*.
- (5) Edwin Green and William Strawderman, Rutgers University: *Bayes Estimates for the Linear Model with Unequal Variances*.
- (b) Maximum Likelihood and Weighted Bootstrapping, Chair: George Casella.
- (1) Adrian Raftery and Michael Newton, University of Washington: *Approximate Bayesian Inference by the Weighted Bootstrap*.
- (2) Charles Geyer, University of Chicago: *Monte Carlo Maximum Likelihood via Gibbs Sampling*.
- (3) Elizabeth Thompson, University of Washington: *Stochastic Simulation for Complex Genetic Analysis*.
- (c) Panel Discussion—Future of Bayesian Inference Using Stochastic Simulation, Chair: Prem Gael.
- Panel—Jim Berger, Alan Gelfand and Adrian Smith.

ACKNOWLEDGMENTS

We are grateful for comments and suggestions from Brad Carlin, Olivier Cappé, David Spiegelhalter, Alan Gelfand, Peter Green, Jun Liu, Sharon McGrayne, Peter Müller, Gareth Roberts and Adrian Smith. Christian Robert's work was partly done during a visit to the Queensland University of Technology, Brisbane, and the author is grateful to Kerrie Mengersen for her hospitality and support. Supported by the Agence Nationale de la Recherche (ANR, 212, rue de Bercy 75012 Paris) through the 2006–2008 project ANR=05-BLAN-0299 Adap'MC. Supported by NSF Grants DMS-06-31632, SES-06-31588 and MMS 1028329.

REFERENCES

- ALBERT, J. H. and CHIB, S. (1993). Bayesian analysis of binary and polychotomous response data. *J. Amer. Statist. Assoc.* **88** 669–679. [MR1224394](#)
- ANDRIEU, C., DE FREITAS, N., DOUCET, A. and JORDAN, M. (2004). An introduction to MCMC for machine learning. *Machine Learning* **50** 5–43.
- ATHREYA, K. B. and NEY, P. (1978). A new approach to the limit theory of recurrent Markov chains. *Trans. Amer. Math. Soc.* **245** 493–501. [MR0511425](#)
- ATHREYA, K. B., DOSS, H. and SETHURAMAN, J. (1996). On the convergence of the Markov chain simulation method. *Ann. Statist.* **24** 69–100. [MR1389881](#)
- BARKER, A. (1965). Monte Carlo calculations of the radial distribution functions for a proton electron plasma. *Aust. J. Physics* **18** 119–133.
- BERTHELSEN, K. and MØLLER, J. (2003). Likelihood and non-parametric Bayesian MCMC inference for spatial point processes based on perfect simulation and path sampling. *Scand. J. Statist.* **30** 549–564.
- BESAG, J. (1974). Spatial interaction and the statistical analysis of lattice systems (with discussion). *J. Roy. Statist. Soc. Ser. B* **36** 192–236. [MR0373208](#)
- BESAG, J. (1975). Statistical analysis of non-lattice data. *The Statistician* **24** 179–195.
- BESAG, J. (1986). On the statistical analysis of dirty pictures. *J. Roy. Statist. Soc. Ser. B* **48** 259–302. [MR0876840](#)
- BESAG, J. and CLIFFORD, P. (1989). Generalized Monte Carlo significance tests. *Biometrika* **76** 633–642. [MR1041408](#)
- BESAG, J., YORK, J. and MOLLIE, A. (1991). Bayesian image restoration, with two applications in spatial statistics (with discussion). *Ann. Inst. Statist. Math.* **43** 1–59. [MR1105822](#)
- BILLINGSLEY, P. (1995). *Probability and Measure*, 3rd ed. Wiley, New York. [MR1324786](#)
- BRONIATOWSKI, M., CELEUX, G. and DIEBOLT, J. (1984). Reconnaissance de mélanges de densités par un algorithme d'apprentissage probabiliste. In *Data Analysis and Informatics, III (Versailles, 1983)* 359–373. North-Holland, Amsterdam. [MR0787647](#)
- BROOKS, S. P., GIUDICI, P. and ROBERTS, G. O. (2003). Efficient construction of reversible jump Markov chain Monte Carlo proposal distributions (with discussion). *J. R. Stat. Soc. Ser. B Stat. Methodol.* **65** 3–55. [MR1959092](#)
- CAPPÉ, O., ROBERT, C. P. and RYDÉN, T. (2003). Reversible jump, birth-and-death and more general continuous time Markov chain Monte Carlo samplers. *J. R. Stat. Soc. Ser. B Stat. Methodol.* **65** 679–700. [MR1998628](#)
- CARLIN, B. and CHIB, S. (1995). Bayesian model choice through Markov chain Monte Carlo. *J. Roy. Statist. Soc. Ser. B* **57** 473–484.
- CARLIN, B., GELFAND, A. and SMITH, A. (1992). Hierarchical Bayesian analysis of change point problems. *Appl. Statist.* **41** 389–405.
- CARPENTER, J., CLIFFORD, P. and FERNHEAD, P. (1997). Building robust simulation-based filters for evolving datasets. Technical report, Dept. Statistics, Oxford Univ.
- CASELLA, G. and GEORGE, E. I. (1992). Explaining the Gibbs sampler. *Amer. Statist.* **46** 167–174. [MR1183069](#)
- CASELLA, G., LAVINE, M. and ROBERT, C. P. (2001). Explaining the perfect sampler. *Amer. Statist.* **55** 299–305. [MR1939363](#)
- CELEUX, G. and DIEBOLT, J. (1985). The SEM algorithm: A probabilistic teacher algorithm derived from the EM algorithm for the mixture problem. *Comput. Statist. Quarter* **2** 73–82.
- CELEUX, G. and DIEBOLT, J. (1990). Une version de type recuit simulé de l'algorithme EM. *C. R. Acad. Sci. Paris Sér. I Math.* **310** 119–124. [MR1044628](#)
- CHAN, K. and GEYER, C. (1994). Discussion of “Markov chains for exploring posterior distribution.” *Ann. Statist.* **22** 1747–1758.

- CHEN, M.-H., SHAO, Q.-M. and IBRAHIM, J. G. (2000). *Monte Carlo Methods in Bayesian Computation*. Springer, New York. [MR1742311](#)
- CHURCHILL, G. (1995). Accurate restoration of DNA sequences (with discussion). In *Case Studies in Bayesian Statistics* (C. Gatsonis and J. S. Hodges, eds.) **2** 90–148. Springer, New York.
- CRESSIE, N. A. C. (1993). *Statistics for Spatial Data*. Wiley, New York. Revised reprint of the 1991 edition. [MR1239641](#)
- DAMIEN, P., WAKEFIELD, J. and WALKER, S. (1999). Gibbs sampling for Bayesian non-conjugate and hierarchical models by using auxiliary variables. *J. R. Stat. Soc. Ser. B Stat. Methodol.* **61** 331–344. [MR1680334](#)
- DEL MORAL, P., DOUCET, A. and JASRA, A. (2006). Sequential Monte Carlo samplers. *J. R. Stat. Soc. Ser. B Stat. Methodol.* **68** 411–436. [MR2278333](#)
- DEMPSTER, A. P., LAIRD, N. M. and RUBIN, D. B. (1977). Maximum likelihood from incomplete data via the EM algorithm (with discussion). *J. Roy. Statist. Soc. Ser. B* **39** 1–38. [MR0501537](#)
- DIEBOLT, J. and ROBERT, C. P. (1994). Estimation of finite mixture distributions through Bayesian sampling. *J. Roy. Statist. Soc. Ser. B* **56** 363–375. [MR1281940](#)
- DIMAKOS, X. K. (2001). A guide to exact simulation. *Internat. Statist. Rev.* **69** 27–48.
- DOUCET, A., DE FREITAS, N. and GORDON, N. (2001). *Sequential Monte Carlo Methods in Practice*. Springer, New York. [MR1847783](#)
- DOUCET, A., GODSILL, S. and ANDRIEU, C. (2000). On sequential Monte Carlo sampling methods for Bayesian filtering. *Statist. Comput.* **10** 197–208.
- DUPUIS, J. A. (1995). Bayesian estimation of movement and survival probabilities from capture–recapture data. *Biometrika* **82** 761–772. [MR1380813](#)
- ECKHARDT, R. (1987). Stan Ulam, John von Neumann, and the Monte Carlo method. *Los Alamos Sci.* **15** Special Issue 131–137. [MR0935772](#)
- EHRMAN, J. R., FOSDICK, L. D. and HANDSCOMB, D. C. (1960). Computation of order parameters in an Ising lattice by the Monte Carlo method. *J. Math. Phys.* **1** 547–558. [MR0117974](#)
- FERNÁNDEZ, R., FERRARI, P. and GARCIA, N. L. (1999). Perfect simulation for interacting point processes, loss networks and Ising models. Technical report, Laboratoire Raphael Salem, Univ. de Rouen.
- FILL, J. A. (1998a). An interruptible algorithm for perfect sampling via Markov chains. *Ann. Appl. Probab.* **8** 131–162. [MR1620346](#)
- FILL, J. (1998b). The move-to front rule: A case study for two perfect sampling algorithms. *Prob. Eng. Info. Sci.* **8** 131–162.
- FISMEN, M. (1998). Exact simulation using Markov chains. Technical Report 6/98, Institutt for Matematiske Fag, Oslo. Diploma-thesis.
- GELFAND, A. E. and DEY, D. K. (1994). Bayesian model choice: Asymptotics and exact calculations. *J. Roy. Statist. Soc. Ser. B* **56** 501–514. [MR1278223](#)
- GELFAND, A. E. and SMITH, A. F. M. (1990). Sampling-based approaches to calculating marginal densities. *J. Amer. Statist. Assoc.* **85** 398–409. [MR1141740](#)
- GELFAND, A. E., SMITH, A. F. M. and LEE, T.-M. (1992). Bayesian analysis of constrained parameter and truncated data problems using Gibbs sampling. *J. Amer. Statist. Assoc.* **87** 523–532. [MR1173816](#)
- GELFAND, A., HILLS, S., RACINE-POON, A. and SMITH, A. (1990). Illustration of Bayesian inference in normal data models using Gibbs sampling. *J. Amer. Statist. Assoc.* **85** 972–982.
- GELMAN, A. and RUBIN, D. (1992). Inference from iterative simulation using multiple sequences (with discussion). *Statist. Sci.* **7** 457–511.
- GEMAN, S. and GEMAN, D. (1984). Stochastic relaxation, Gibbs distributions and the Bayesian restoration of images. *IEEE Trans. Pattern Anal. Mach. Intell.* **6** 721–741.
- GEORGE, E. and MCCULLOCH, R. (1993). Variable selection via Gibbs sampling. *J. Amer. Statist. Assoc.* **88** 881–889.
- GEORGE, E. I. and ROBERT, C. P. (1992). Capture–recapture estimation via Gibbs sampling. *Biometrika* **79** 677–683. [MR1209469](#)
- GEYER, C. (1992). Practical Monte Carlo Markov chain (with discussion). *Statist. Sci.* **7** 473–511.
- GEYER, C. J. and MØLLER, J. (1994). Simulation procedures and likelihood inference for spatial point processes. *Scand. J. Statist.* **21** 359–373. [MR1310082](#)
- GEYER, C. and THOMPSON, E. (1995). Annealing Markov chain Monte Carlo with applications to ancestral inference. *J. Amer. Statist. Assoc.* **90** 909–920.
- GILKS, W. (1992). Derivative-free adaptive rejection sampling for Gibbs sampling. In *Bayesian Statistics 4* (J. M. Bernardo, J. O. Berger, A. P. Dawid and A. F. M. Smith, eds.) 641–649. Oxford Univ. Press, Oxford.
- GILKS, W. R. and BERZUINI, C. (2001). Following a moving target—Monte Carlo inference for dynamic Bayesian models. *J. R. Stat. Soc. Ser. B Stat. Methodol.* **63** 127–146. [MR1811995](#)
- GILKS, W., BEST, N. and TAN, K. (1995). Adaptive rejection Metropolis sampling within Gibbs sampling. *Appl. Statist.* **44** 455–472.
- GILKS, W. R., ROBERTS, G. O. and SAHU, S. K. (1998). Adaptive Markov chain Monte Carlo through regeneration. *J. Amer. Statist. Assoc.* **93** 1045–1054. [MR1649199](#)
- GORDON, N., SALMOND, J. and SMITH, A. (1993). A novel approach to non-linear/non-Gaussian Bayesian state estimation. *IEEE Proceedings on Radar and Signal Processing* **140** 107–113.
- GREEN, P. J. (1995). Reversible jump Markov chain Monte Carlo computation and Bayesian model determination. *Biometrika* **82** 711–732. [MR1380810](#)
- GUIHENNEUC-JOUYAU, C. and ROBERT, C. P. (1998). Discretization of continuous Markov chains and Markov chain Monte Carlo convergence assessment. *J. Amer. Statist. Assoc.* **93** 1055–1067. [MR1649200](#)
- HAMMERSLEY, J. M. (1974). Discussion of Mr Besag’s paper. *J. Roy. Statist. Soc. Ser. B* **36** 230–231.
- HAMMERSLEY, J. M. and HANDSCOMB, D. C. (1964). *Monte Carlo Methods*. Wiley, New York.
- HAMMERSLEY, J. M. and MORTON, K. W. (1954). Poor man’s Monte Carlo. *J. Roy. Statist. Soc. Ser. B* **16** 23–38. [MR0064475](#)
- HANDSCHIN, J. E. and MAYNE, D. Q. (1969). Monte Carlo techniques to estimate the conditional expectation in multi-stage non-linear filtering. *Internat. J. Control* **9** 547–559. [MR0246490](#)

- HASTINGS, W. (1970). Monte Carlo sampling methods using Markov chains and their application. *Biometrika* **57** 97–109.
- HITCHCOCK, D. B. (2003). A history of the Metropolis–Hastings algorithm. *Amer. Statist.* **57** 254–257. [MR2037852](#)
- HOBERT, J. P. and CASELLA, G. (1996). The effect of improper priors on Gibbs sampling in hierarchical linear mixed models. *J. Amer. Statist. Assoc.* **91** 1461–1473. [MR1439086](#)
- HOBERT, J. P. and MARCHEV, D. (2008). A theoretical comparison of the data augmentation, marginal augmentation and PX-DA algorithms. *Ann. Statist.* **36** 532–554. [MR2396806](#)
- HOBERT, J. P., JONES, G. L., PRESNELL, B. and ROSENTHAL, J. S. (2002). On the applicability of regenerative simulation in Markov chain Monte Carlo. *Biometrika* **89** 731–743. [MR1946508](#)
- JONES, G. L. and HOBERT, J. P. (2001). Honest exploration of intractable probability distributions via Markov chain Monte Carlo. *Statist. Sci.* **16** 312–334. [MR1888447](#)
- JONES, G. L., HARAN, M., CAFFO, B. S. and NEATH, R. (2006). Fixed-width output analysis for Markov chain Monte Carlo. *J. Amer. Statist. Assoc.* **101** 1537–1547. [MR2279478](#)
- KEMENY, J. G. and SNELL, J. L. (1960). *Finite Markov Chains*. Van Nostrand, Princeton, NJ. [MR0115196](#)
- KENDALL, W. S. and MØLLER, J. (2000). Perfect simulation using dominating processes on ordered spaces, with application to locally stable point processes. *Adv. in Appl. Probab.* **32** 844–865. [MR1788098](#)
- KIPNIS, C. and VARADHAN, S. R. S. (1986). Central limit theorem for additive functionals of reversible Markov processes and applications to simple exclusions. *Comm. Math. Phys.* **104** 1–19. [MR0834478](#)
- KIRKPATRICK, S., GELATT, C. D. JR. and VECCHI, M. P. (1983). Optimization by simulated annealing. *Science* **220** 671–680. [MR0702485](#)
- KITAGAWA, G. (1996). Monte Carlo filter and smoother for non-Gaussian nonlinear state space models. *J. Comput. Graph. Statist.* **5** 1–25. [MR1380850](#)
- KONG, A., LIU, J. and WONG, W. (1994). Sequential imputations and Bayesian missing data problems. *J. Amer. Statist. Assoc.* **89** 278–288.
- KUHN, T. (1996). *The Structure of Scientific Revolutions*, 3rd ed. Univ. Chicago Press, Chicago.
- LANDAU, D. P. and BINDER, K. (2005). *A Guide to Monte Carlo Simulations in Statistical Physics*. Cambridge Univ. Press, Cambridge.
- LANGE, N., CARLIN, B. P. and GELFAND, A. E. (1992). Hierarchical Bayes models for the progression of HIV infection using longitudinal CD4 T-cell numbers. *J. Amer. Statist. Assoc.* **87** 615–626.
- LAWRENCE, C. E., ALTSCHUL, S. F., BOGUSKI, M. S., LIU, J. S., NEUWALD, A. F. and WOOTTON, J. C. (1993). Detecting subtle sequence signals: A Gibbs sampling strategy for multiple alignment. *Science* **262** 208–214.
- LIU, J. and CHEN, R. (1995). Blind deconvolution via sequential imputations. *J. Amer. Statist. Assoc.* **90** 567–576.
- LIU, C. and RUBIN, D. B. (1994). The ECME algorithm: A simple extension of EM and ECM with faster monotone convergence. *Biometrika* **81** 633–648. [MR1326414](#)
- LIU, J. S. and WU, Y. N. (1999). Parameter expansion for data augmentation. *J. Amer. Statist. Assoc.* **94** 1264–1274. [MR1731488](#)
- LIU, J. S., WONG, W. H. and KONG, A. (1994). Covariance structure of the Gibbs sampler with applications to the comparisons of estimators and augmentation schemes. *Biometrika* **81** 27–40. [MR1279653](#)
- LIU, J. S., WONG, W. H. and KONG, A. (1995). Covariance structure and convergence rate of the Gibbs sampler with various scans. *J. Roy. Statist. Soc. Ser. B* **57** 157–169. [MR1325382](#)
- MADRAS, N. and SLADE, G. (1993). *The Self-Avoiding Walk*. Birkhäuser, Boston, MA. [MR1197356](#)
- MARIN, J.-M. and ROBERT, C. P. (2007). *Bayesian Core: A Practical Approach to Computational Bayesian Statistics*. Springer, New York. [MR2289769](#)
- MARSHALL, A. (1965). The use of multi-stage sampling schemes in Monte Carlo computations. In *Symposium on Monte Carlo Methods*. Wiley, New York.
- MENG, X.-L. and RUBIN, D. B. (1993). Maximum likelihood estimation via the ECM algorithm: A general framework. *Biometrika* **80** 267–278. [MR1243503](#)
- MENG, X.-L. and VAN DYK, D. A. (1999). Seeking efficient data augmentation schemes via conditional and marginal augmentation. *Biometrika* **86** 301–320. [MR1705351](#)
- MENGERSEN, K. L. and TWEEDIE, R. L. (1996). Rates of convergence of the Hastings and Metropolis algorithms. *Ann. Statist.* **24** 101–121. [MR1389882](#)
- METROPOLIS, N. (1987). The beginning of the Monte Carlo method. *Los Alamos Sci.* **15** Special Issue 125–130. [MR0935771](#)
- METROPOLIS, N. and ULAM, S. (1949). The Monte Carlo method. *J. Amer. Statist. Assoc.* **44** 335–341. [MR0031341](#)
- METROPOLIS, N., ROSENBLUTH, A., ROSENBLUTH, M., TELLER, A. and TELLER, E. (1953). Equations of state calculations by fast computing machines. *J. Chem. Phys.* **21** 1087–1092.
- MØLLER, J. and WAAGEPETERSEN, R. (2003). *Statistical Inference and Simulation for Spatial Point Processes*. Chapman & Hall/CRC, Boca Raton.
- MOUSSOURIS, J. (1974). Gibbs and Markov random systems with constraints. *J. Statist. Phys.* **10** 11–33. [MR0432132](#)
- MYKLAND, P., TIERNEY, L. and YU, B. (1995). Regeneration in Markov chain samplers. *J. Amer. Statist. Assoc.* **90** 233–241. [MR1325131](#)
- NEAL, R. (1996). Sampling from multimodal distributions using tempered transitions. *Statist. Comput.* **6** 353–356.
- NEAL, R. M. (2003). Slice sampling (with discussion). *Ann. Statist.* **31** 705–767. [MR1994729](#)
- PEARL, J. (1987). Evidential reasoning using stochastic simulation of causal models. *Artificial Intelligence* **32** 245–257. [MR0885357](#)
- PESKUN, P. H. (1973). Optimum Monte-Carlo sampling using Markov chains. *Biometrika* **60** 607–612. [MR0362823](#)
- PESKUN, P. H. (1981). Guidelines for choosing the transition matrix in Monte Carlo methods using Markov chains. *J. Comput. Phys.* **40** 327–344. [MR0617102](#)
- PROPP, J. G. and WILSON, D. B. (1996). Exact sampling with coupled Markov chains and applications to statistical mechanics. In *Proceedings of the Seventh International Conference on Random Structures and Algorithms (Atlanta, GA, 1995)* **9** 223–252. [MR1611693](#)
- QIAN, W. and TITTERINGTON, D. M. (1990). Parameter estimation for hidden Gibbs chains. *Statist. Probab. Lett.* **10** 49–58. [MR1060302](#)

- RAFTERY, A. and BANFIELD, J. (1991). Stopping the Gibbs sampler, the use of morphology, and other issues in spatial statistics. *Ann. Inst. Statist. Math.* **43** 32–43.
- RICHARDSON, S. and GREEN, P. J. (1997). On Bayesian analysis of mixtures with an unknown number of components (with discussion). *J. Roy. Statist. Soc. Ser. B* **59** 731–792. [MR1483213](#)
- RIPLEY, B. D. (1987). *Stochastic Simulation*. Wiley, New York. [MR0875224](#)
- ROBERT, C. P. (1995). Convergence control methods for Markov chain Monte Carlo algorithms. *Statist. Sci.* **10** 231–253. [MR1390517](#)
- ROBERT, C. P. and CASELLA, G. (2004). *Monte Carlo Statistical Methods*, 2nd ed. Springer, New York. [MR2080278](#)
- ROBERTS, G. O. and ROSENTHAL, J. S. (1999). Convergence of slice sampler Markov chains. *J. R. Stat. Soc. Ser. B Stat. Methodol.* **61** 643–660. [MR1707866](#)
- ROBERTS, G. O. and ROSENTHAL, J. S. (2007). Coupling and ergodicity of adaptive Markov chain Monte Carlo algorithms. *J. Appl. Probab.* **44** 458–475. [MR2340211](#)
- ROBERTS, G. O., GELMAN, A. and GILKS, W. R. (1997). Weak convergence and optimal scaling of random walk Metropolis algorithms. *Ann. Appl. Probab.* **7** 110–120. [MR1428751](#)
- ROSENBLUTH, M. and ROSENBLUTH, A. (1955). Monte Carlo calculation of the average extension of molecular chains. *J. Chem. Phys.* **23** 356–359.
- ROSENTHAL, J. S. (1995). Minorization conditions and convergence rates for Markov chain Monte Carlo. *J. Amer. Statist. Assoc.* **90** 558–566. [MR1340509](#)
- RUBIN, D. (1978). Multiple imputation in sample surveys: A phenomenological Bayesian approach to nonresponse. In *Imputation and Editing of Faulty or Missing Survey Data*. U.S. Department of Commerce, Washington, DC.
- SMITH, A. F. M. and GELFAND, A. E. (1992). Bayesian statistics without tears: A sampling-resampling perspective. *Amer. Statist.* **46** 84–88. [MR1165566](#)
- STEPHENS, D. A. (1994). Bayesian retrospective multiple-change point identification. *Appl. Statist.* **43** 159–178.
- STEPHENS, M. (2000). Bayesian analysis of mixture models with an unknown number of components—an alternative to reversible jump methods. *Ann. Statist.* **28** 40–74. [MR1762903](#)
- STEPHENS, D. A. and SMITH, A. F. M. (1993). Bayesian inference in multipoint gene mapping. *Ann. Hum. Genetics* **57** 65–82.
- TANNER, M. A. and WONG, W. H. (1987). The calculation of posterior distributions by data augmentation (with discussion). *J. Amer. Statist. Assoc.* **82** 528–550. [MR0898357](#)
- TIERNEY, L. (1994). Markov chains for exploring posterior distributions (with discussion). *Ann. Statist.* **22** 1701–1786. [MR1329166](#)
- TITTERINGTON, D. and COX, D. R. (2001). *Biometrika: One Hundred Years*. Oxford Univ. Press, Oxford. [MR1841263](#)
- WAKEFIELD, J., SMITH, A., RACINE-POON, A. and GELFAND, A. (1994). Bayesian analysis of linear and non-linear population models using the Gibbs sampler. *Appl. Statist.* **43** 201–222.
- WANG, C. S., RUTLEDGE, J. J. and GIANOLA, D. (1993). Marginal inferences about variance-components in a mixed linear model using Gibbs sampling. *Gen. Sel. Evol.* **25** 41–62.
- WANG, C. S., RUTLEDGE, J. J. and GIANOLA, D. (1994). Bayesian analysis of mixed linear models via Gibbs sampling with an application to litter size in Iberian pigs. *Gen. Sel. Evol.* **26** 91–115.
- WEI, G. and TANNER, M. (1990). A Monte Carlo implementation of the EM algorithm and the poor man's data augmentation algorithm. *J. Amer. Statist. Assoc.* **85** 699–704.
- ZEGER, S. L. and KARIM, M. R. (1991). Generalized linear models with random effects; a Gibbs sampling approach. *J. Amer. Statist. Assoc.* **86** 79–86. [MR1137101](#)